

## Suffix trees

### Data structures for string pattern matching: Suffix trees

- Linear algorithms for exact string matching
  - KMP
  - Z-value algorithm
- What is suffix tree?
  - A tree-like **data structure** for solving problems involving strings.
  - Related data structures: Trie (**re**trieval) & PATRICIA (radix tree)
  - Allow the storage of **all substrings** of a given string in linear space
  - Simple algorithm to solve string pattern matching problem in linear time

## Better than hash tables?

- Hash tables are certainly easier to understand. And, one can produce a hash table of all length  $k$  strings in  $O(m)$  time and look up a  $k$ -length string  $x$  in  $O(k)$  time, finding all  $p$  places where string  $x$  is found in  $O(p)$  time. This is the same as the bound for suffix trees.
- What if you don't know how long the string  $x$  is going to be?
- And most other string matching tricks don't work for it either.

## Suffix Tree: definition

- A suffix tree  $ST$  for an  $m$ -character string  $S$  is a rooted directed tree with exactly  $m$  leaves numbered 1 to  $m$ .
- Each internal node, other than the root, has at least two children and each edge is labeled with a nonempty substring of  $S$ .

## Suffix tree: definition

- No two edges out of a node can have edge-labels beginning with the same character.
- The key feature of the suffix tree is that for any leaf  $i$ , the concatenation of the edge-labels on the path from the root to the leaf  $i$  exactly spells out the suffix of  $S$  that starts at position  $i$ .

## Suffix Trees: Example


- Suffixes of 'papua'
  - 'papua'
  - 'apua'
  - 'pua'
  - 'ua'
  - 'a'
  - ''

## Suffix Trees: Example

- Suffixes of 'papua'
    - 'papua'
    - 'apua'
    - 'pua'
    - 'ua'
    - 'a'
    - ''
- NOTE: Assume the string terminates with some character found nowhere else in the string. (eg. '\0')


## Suffix Trees: Example

- Suffix tree for 'papua'




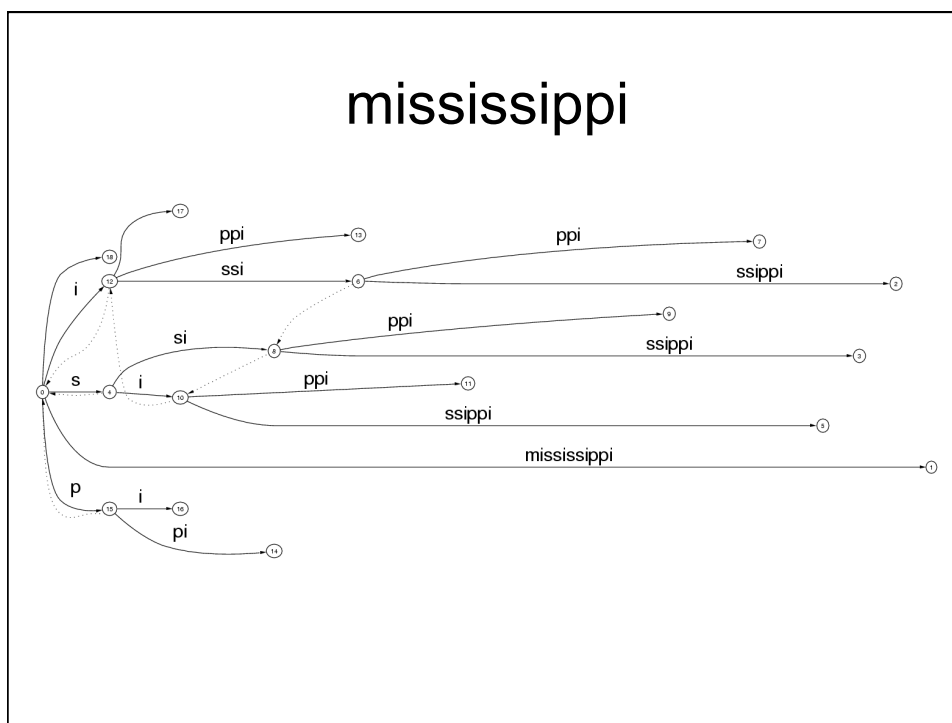
## Suffix Trees: Example

- Suffix tree for 'papua'



## caracas





# Suffix Trees

- Exact matching in linear time
- Many others
- “We know of no other single data structure that allows efficient solutions to such a wide range of complex string problems.” - Dan Gusfield

## Exact string matching problem

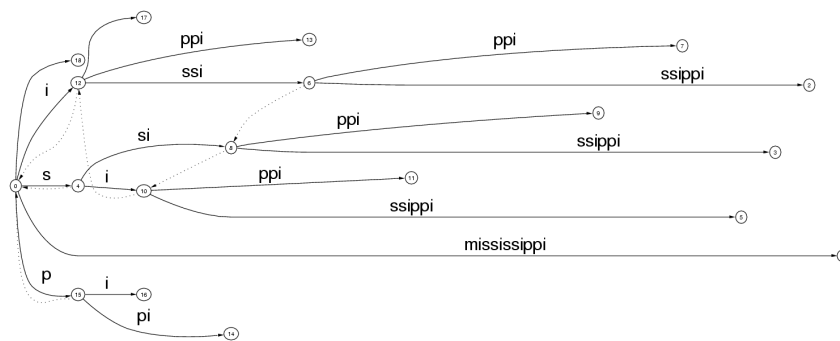
- Given a pattern  $P$  of length  $m$ , find all occurrences of  $P$  in text  $T$ 
  - $O(n+m)$  algorithm
- Solution: Build a suffix tree  $ST$  for text  $T$  in  $O(m)$  time. Then, match the characters of  $P$  along the unique path in  $ST$  until either  $P$  is exhausted or no more matches are possible.

## Exact string matching problem


- Find 'ssi' in 'mississippi'

## Exact string matching problem

- Find 'ssi' in 'mississippi'

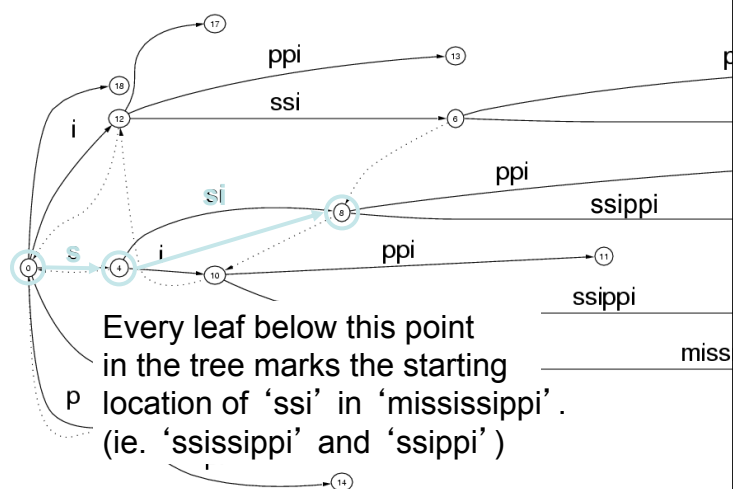


## Exact string matching problem






## Exact string matching problem




## Exact string matching problem

- Find 'sissy' in 'mississippi'

## Exact string matching problem



# Exact string matching problem



## Comparing to the other algorithms

- KMP and Boyer-Moore both achieve this worst case bound.
  - $O(m+n)$  when the text and pattern are presented together.
- Suffix trees are much faster when the text is fixed and known first while the patterns vary.
  - $O(m)$  for single time processing the text, then only  $O(n)$  for each new pattern.
- Based on suffix trees, is faster for searching a number of patterns at one time against a single text (**exact set matching problem**)
  - Aho-Corasick algorithm: preprocessing P instead of T.

## Building the Suffix Tree

- How do we build a suffix tree?  
`while suffixes remain:`  
`add next shortest suffix to the tree`


## Building the Suffix Tree

- papua




## Building the Suffix Tree

- papua




## Building the Suffix Tree

- papua




## Building the Suffix Tree

- papua




## Building the Suffix Tree

- papua




## Building the Suffix Tree

- papua



## Building the Suffix Tree

- papua



## Building the Suffix Tree

- How do we build a suffix tree?

```
while suffices remain:
    add next shortest suffix to the tree
```

Naïve method -  $O(m^2)$  ( $m$  = text size)

## Building the Suffix Tree in $O(m)$ time

- In the previous example, we assumed that the tree can be built in  $O(m)$  time.
- Weiner showed original  $O(m)$  algorithm (Knuth is claimed to have called it “the algorithm of 1973”)
- More space efficient algorithm by McCreight in 1976
- Simpler ‘on-line’ algorithm by Ukkonen in 1995


## Ukkonen’s Algorithm

- Build suffix tree  $T$  for string  $S[1..m]$ 
  - Build the tree in  $m$  phases, one for each character. At the end of phase  $i$ , we will have tree  $T_i$ , which is the tree representing the prefix  $S[1..i]$  → online construction
    - In each phase  $i$ , we have  $i$  extensions, one for each character in the current prefix. At the end of extension  $j$ , we will have ensured that  $S[j..i]$  is in the tree  $T_i$ .




## Implicit suffix tree

- papua




## Implicit suffix tree

- papua



## Implicit suffix tree

- papua



Implicit suffix tree can be transformed from/into regular suffix tree in  $O(n)$  time.

## Ukkonen's Algorithm


Pseudo code for **Ukk**:

```

Construct tree  $T_1$ .
for  $i = 1$  to  $m-1$  do
  begin {phase  $i+1$ }
    for  $j = 1$  to  $i+1$  do
      begin {extension  $j$ }
        In the current tree find the end of the path from the root
        labeled  $t[j \dots i]$ . If necessary, extend that path by adding
        character  $t[i+1]$ , thus ensuring that string  $t[j \dots i+1]$  is in the
        tree.
      end;
    end;
  end;
end;
```

## An Exmample

$t = a c c a \$$



## Ukkonen's Algorithm


- This is an  $O(m^3)$  time,  $O(m^2)$  space algorithm.
- We need a few implementation speed-ups to achieve the  $O(m)$  time and  $O(m)$  space bounds.

## Suffix extension rules

- 3 possible ways to extend  $S[j..i]$  with character  $i+1$ .
  - $S[j..i]$  ends at a leaf. Add the character  $i+1$  to the end of the leaf edge.
  - No path from the end of  $S[j..i]$  starts with the  $i+1$  character. Split the edge and create a new node if necessary, then add a new leaf with character  $i+1$ .  
(This is the only extension that increases the number of leaves! The new leaf represents the suffix starting at position  $j$ .)
  - There is already a path from the end of  $S[j..i]$  starts with the  $i+1$  character, or  $S[j..i+1]$  correspond to a path. Do nothing.


## Ukkonen's Algorithms

$t = a c c a \$$   
 $t[1..3] = acc$   
 $t[1..4] = acca$




## Ukkonen's Algorithm: Speed-up 1

- Suffix Links
  - speed up navigation to the next extension point in the tree




## Ukkonen's Algorithm: Speed-up 2

- Skip/Count Trick
  - instead of stepping through each character, we know that we can just jump, as long as we're the right distance




## Ukkonen's Algorithm: Speed-up 2

- Skip/Count Trick
  - instead of stepping through each character, we know that we can just jump, as long as we're the right distance




## Ukkonen's Algorithm: Speed-up 3

- Edge-Label Compression
  - since we have a copy of the string, we don't need to store copies of the substrings for each edge



## Ukkonen's Algorithm: Speed-up 3

- Edge-Label Compression
  - since we have a copy of the string, we don't need to store copies of the substrings for each edge
  - $O(m^2)$  space becomes  $O(m)$  space



## Ukkonen's Algorithm: Speed-up 4

- A match is a show stopper.
  - If we find a match to our next character (rule 3 applies), we're done this phase.

## Ukkonen's Algorithm: Speed-up 5

- Once a leaf, always a leaf (implicitly implement rule 1).
  - We don't need to update each leaf, since it will always be the end of the current string. We can get these updates for free.
  - Either 1) maintain a global *end-of-string* index or 2) insert the whole string for every leaf

## Ukkonen's Algorithm: Speed-ups

- Because of speed-ups 4 and 5, we can pick up the next phase right where we ended the last one!



## Ukkonen's Algorithm – mississippi with Speed-ups

```
void SuffixTree::update(char* s, int len) {
    ...
    int i;
    int j;
    ...
    for (i = 0, j = 0; i < len; i++) {
        while (j <= i) {
            ... all the work ...
        }
    }
}
```

## Possible execution

### Extension Case distribution

j=	1	2	3	4	5	6	7	8	9=m	
i=1:	1	3								Need to compute:
i=2:	1	1	2							At most one Case 2 per i
i=3:	1	1	2	2						At most one Case 1 or 3 per j
i=4:	1	1	3	2	2					
i=5:	1	1	1	2	2	2				
i=6:	1	1	1	3	3	2	2			
i=7:	1	1	1	1	1	3	2	2		
i=8:	1	1	1	1	1	1	3	2	2	

## Creating a true suffix tree

- Run another iteration of Ukkonen algorithm on  $S\$$
- No suffix is now a prefix of any other suffix.
- As a result, each suffix will end at a leaf.
- Replace each index on every leaf edge with the number  $m$ .

**Total Algorithm time  $O(m)$**


## Ukkonen's Algorithm – The Punch Line

- By combining all of the speed-ups, we can now construct a suffix tree  $T_m$  representing the string  $S[1..m]$  in
  - $O(m)$  time and in
  - $O(m)$  space!

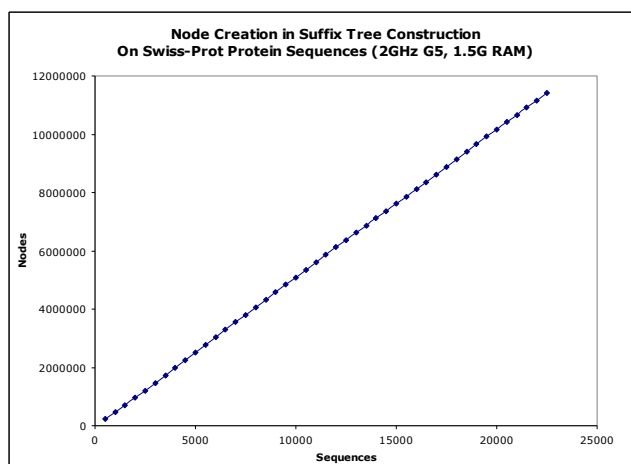
## Exact string matching

- Both  $P$  ( $|P|=n$ ) and  $T$  ( $|T|=m$ ) are known:
  - Suffix tree method achieves same worst-case bound  $O(n+m)$  as KMP.
- $T$  is fixed and build suffix tree, then  $P$  is input,  $k$  is the number of occurrences of  $P$ 
  - Using suffix tree:  $O(n+k)$
  - In contrast (KMP, preprocess  $P$ ):  $O(n+m)$  for any single  $P$
- $P$  is fixed, then  $T$  is input
  - Selecting KMP rather than suffix tree
  - or Aho-Corasick algorithm (exact set matching problem)

## Ukkonen's Algorithm – Time Performance



## Ukkonen's Algorithm – Memory Usage



## Applications

- Problems
  - linear-time longest common substring
  - constant-time least common ancestor
  - maximally repetitive structures
  - all-pairs suffix-prefix matching
  - compression
  - inexact matching
  - conversion to suffix arrays

## Bioinformatics applications

- Applications
  - Sequence comparison
  - motif discovery
  - PST – probabilistic suffix trees
  - SVM string kernels
  - chromosome-level similarities and rearrangements